AEREALITH AI

Market Intelligence & Launch Strategy

A decision-grade review of demand, competitors, pricing, market size, positioning, trust requirements, and go-to-market sequencing.

FULL MARKET RESEARCH

SOURCES + MODELED ESTIMATES

Prepared 13 August 2026

Confidential strategic planning document | Version 1.0

A companion for a more connected world.

The market is large. The winnable entry point is deliberately narrow.

Aerealith should launch as the permission-aware AI operations layer for technical communities and developer-led organizations - not as another universal chatbot.

88%
ORGANIZATIONS USING AI
Survey-reported use in at least one function, 2025.

23%
SCALING AI AGENTS
Only a minority report scaled agent deployments.

46%
DEVELOPERS DISTRUST ACCURACY
Trust is weaker than adoption among developers.

$460M
MODELED 5-YEAR SAM
Bottom-up planning estimate; not a published forecast.

INVESTMENT THESIS

Aerealith can win where assistants, organizational memory, community operations, and infrastructure awareness intersect. The launch product must make source permissions, memory provenance, and every consequential action visible. That control plane is more defensible than a broad feature list and more credible than leading with speculative companion claims.

01 1. Win the beachhead

Technical community operators already live across Discord, GitHub, calendars, documents, and infrastructure. Their context is fragmented, their automations are brittle, and incumbent copilots do not own the community-to-operations workflow.

02 2. Prove the control layer

Ship cited answers, memory receipts, permission scopes, an action preview, approval gates, and a durable action ledger before promoting autonomous workflows. Trust is part of the product surface, not a compliance appendix.

03 3. Expand by adjacency

Move from individuals and guilds to developer-led teams, then to hybrid and self-hosted business deployments. Preserve the same identity, context, and governance primitives at every stage.

What this report measures - and what it refuses to pretend is known

The analysis triangulates public adoption data, first-party competitor materials, internal Aerealith plans, and a transparent bottom-up model. Prices and product claims are dated 13 August 2026.

Evidence hierarchy

PUBLISHED DATA
Government, intergovernmental, and independent survey material.

VENDOR CLAIM
Pricing, feature, user-count, and adoption statements from product owners.

MODELED ESTIMATE
Aerealith TAM, SAM, SOM, pricing, and unit-economics scenarios built in this report.

Question	Method	Decision supported
Is demand real?	Adoption surveys + usage proxies	Category timing
Who buys first?	Pain/urgency/fit scoring	Beachhead ICP
Who already owns the job?	Competitor category map + pricing	Positioning
How large is the target?	Bottom-up accounts x realistic ARR	Capital allocation
What must ship?	Gap analysis vs trust and workflow needs	Roadmap
How can it reach buyers?	Channel economics + funnel design	GTM

Research boundaries

In scope
Governed AI assistants, enterprise knowledge, workflow automation, personal memory, community operations, developer and infrastructure workflows, cloud/local/hybrid delivery.

Out of scope as market size
Total AI infrastructure spend, all SaaS, all small-business software, consumer entertainment AI, and generalized humanoid or AGI forecasts.

Limitation
Private-company revenue, retention, and product quality are rarely public. Feature scoring reflects documented capability, not hands-on benchmark testing.

INTERPRETATION RULES

- Survey percentages are not directly interchangeable; bases, geographies, and definitions differ.
- A vendor's feature page proves positioning and documented capability - not adoption, reliability, or customer satisfaction.
- The TAM model is a planning boundary. The SOM scenarios are outcomes to earn, not forecasts to repeat as fact.
- Aerealith capability is labeled current-plan or target-state; it is never scored as already shipped.
- Every material numeric claim maps to a source ID or is explicitly marked modeled.

Aerealith already has a coherent product idea - but not yet a coherent market promise

The internal plan combines personal assistance, memory, automation, community operations, infrastructure visibility, analytics, extensibility, and long-horizon research.

Assistant

Conversational help, identities, tone, summaries, tasks and reminders.

Memory

User-visible memories, context retrieval, retention controls and dashboard concepts.

Automation

Triggers, actions, workflows and approval-gated execution.

Community

Discord-first coordination, moderation-adjacent support and shared context.

Developer + infra

GitHub, read-only Proxmox and infrastructure-health awareness.

Platform

Integrations, API/SDK, local runtime, hybrid deployment and future marketplace.

Internal decision area	Brand guide	Master plan	Pricing v3	Market implication
Primary audience	Broad connected world	Power users, devs, homelab, creators	Individuals through enterprise	Launch audience must narrow
Hosting	Local-first bridge / hybrid	Helix-hosted; no launch self-host	Open-source self-host runtime	Resolve promise and sequence
Public tiers	Not specified	$0 / $5 / $10 / $20 / Enterprise	$0 / $3 / $6 / $10 / $15 / $20 / $25	One public ladder is required
Commercial lead	Secure command center	Personal + infrastructure assistant	Feature/usage value	Lead with a concrete job

Finding: this is a positioning and sequencing problem, not a lack-of-vision problem. The product thesis becomes credible when the first buyer, first workflow, and first trust promise are explicit.

Recommended beachhead: technical communities that operate like small organizations

Discord communities, open-source projects, creator teams, game studios, and homelab groups share a distinctive problem: conversation, code, knowledge, and systems live in different places.

The urgent job

Answer questions with permission-aware context from Discord, GitHub, documents, calendars, and infrastructure - then preview and execute approved actions with a visible ledger.

Why incumbents leave room

Suite copilots are strongest inside their own ecosystems. Automation tools expose power but demand builder effort. Community bots solve narrow moderation jobs. Aerealith can connect the seams.

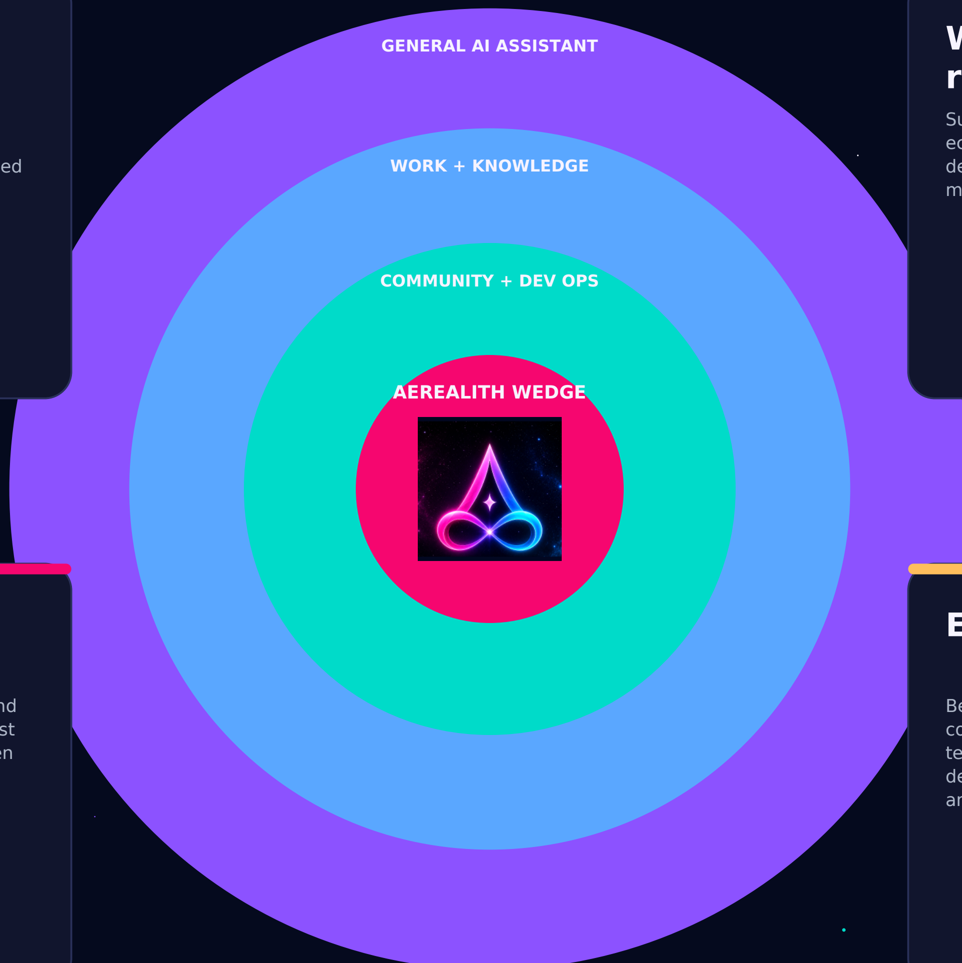

Proof before expansion

Activation must demonstrate one cited answer and one approved cross-system action. Retention must come from durable context, not novelty. Only then should broader personal-companion messaging expand.

Economic buyer

Begin with the person who already bears coordination cost: guild owner, project maintainer, technical community manager, studio ops lead, or developer-founder. Their pain is measurable in time and incidents.

Aerealith competes across four categories - while selling one outcome

The product cannot be evaluated inside a single analyst category. Buyers will compare it with the tool already closest to their daily pain.

AI ASSISTANT

Chat, reasoning, drafting, personal productivity

OpenAI | Claude | Gemini | Copilot

CONTEXT + KNOWLEDGE

Permissions, search, memory, citations, organizational graph

Glean | Guru | Notion | Rovo

ACTION + AUTOMATION

Triggers, workflows, integrations, approvals, execution

Zapier | n8n | Make | Dify

COMMUNITY + OPERATIONS

Discord coordination, support, moderation, dev/infra awareness

MEE6 | Dyno | Carl-bot | Ticket Tool

- Category rule: market against the incumbent that owns the buyer's current workaround, then explain why Aerealith connects context and action without surrendering control.

Adoption is broad; operational maturity is not

This is favorable for a control-oriented product. Buyers have experienced AI's value, but many have not solved integration, governance, reliability, or workflow redesign.

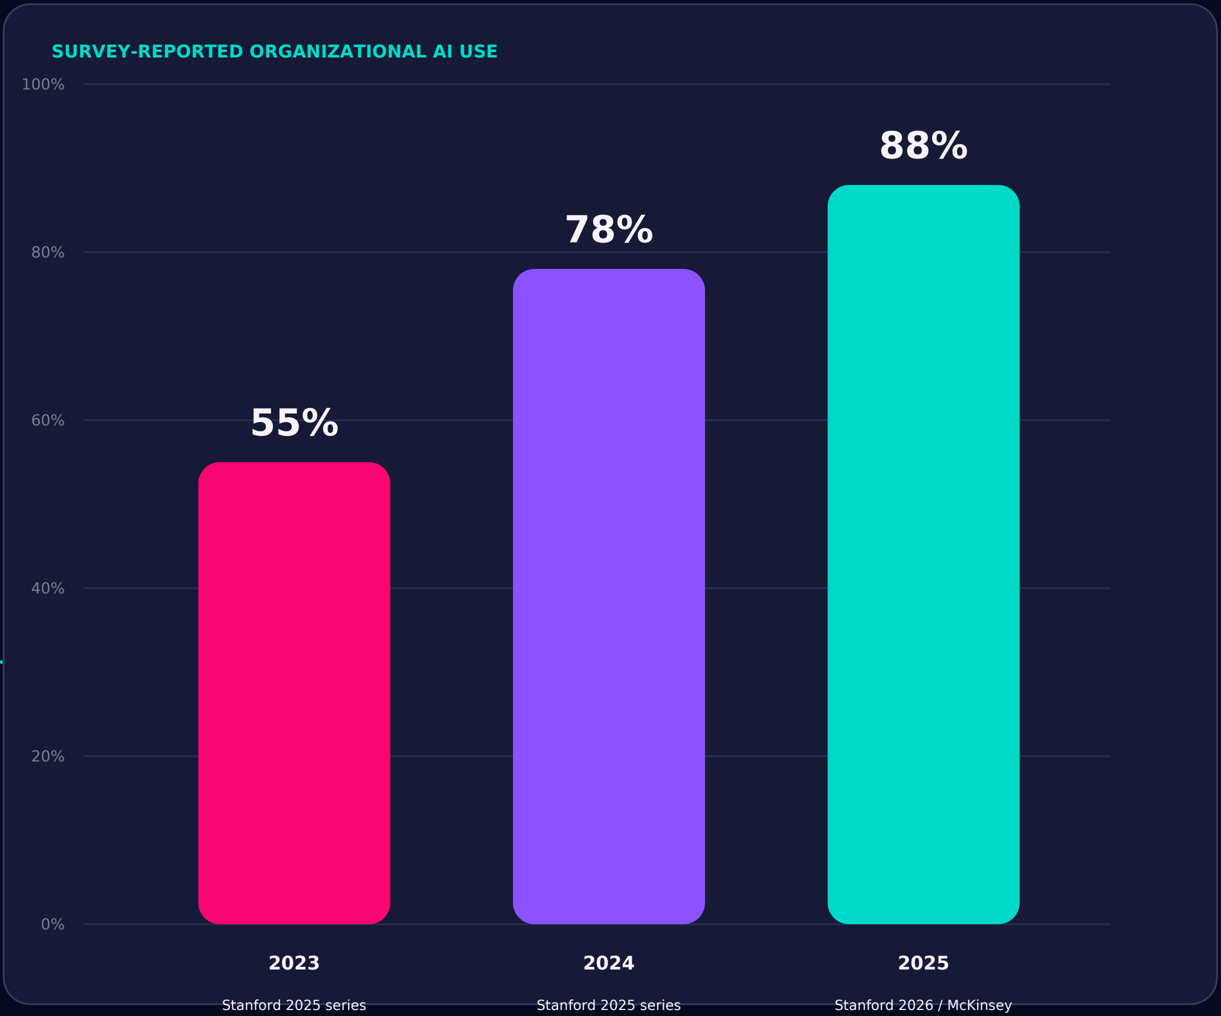

62%

EXPERIMENTING WITH OR SCALING AGENTS

McKinsey 2025; 23% report scaling.

17-20%

U.S. FIRMS USING AI

Census range, Dec 2025-May 2026.

20.2%

OECD FIRMS USING AI

Up from 8.7% in 2023.

29%

GENAI-USING SMEs APPLY IT TO CORE WORK

Most SME use remains peripheral or simple.

The opportunity is the gap between trying AI and trusting it with work

Aerealith should monetize operational confidence: context integrity, scoped permissions, observable automation, and recoverable actions.

ILLUSTRATIVE MATURITY FUNNEL

1 **TRY** Chat and draft

2 **CONNECT** Bring internal context

3 **STANDARDIZE** Shared workflows

4 **DELEGATE** Agents take actions

5 **GOVERN** Measure and improve

Percentages anchor the first and delegation stages; intermediate widths are analytical, not measured conversion rates.

Friction 1: context

Answers are only useful when the source is current, permitted, cited, and understandable.

Friction 2: consequence

Every action introduces error, security, and accountability risk. Preview, approval, rollback, and logs reduce that risk.

Friction 3: ownership

Teams need to know who configured the agent, who can override it, where data flows, and what happens when a model fails.

Product thesis: the trust surface is the adoption engine. 'Permission-aware' must be demonstrated in onboarding and daily use, not buried in enterprise settings.

Six forces are expanding the category - and six forces can stall it

Aerealith's launch plan should amplify structural tailwinds while designing around trust, cost, distribution, and incumbent bundling.

TAILWINDS

- Model and platform spending is expanding rapidly, increasing capability and buyer awareness.
- Agents are becoming expected features inside enterprise applications.
- Developer ecosystems and AI-enabled repositories are growing quickly.
- Organizations have experienced assistant value and now seek integration into real workflows.
- Discord and developer platforms provide native app distribution and monetization paths.
- Open models and local runtimes make hybrid delivery strategically possible.

HEADWINDS

- Incumbents bundle assistants into suites, compressing willingness to pay for generic chat.
- Accuracy distrust remains high among developers - a core launch audience.
- Agent actions introduce security, privacy, and accountability risk.
- Inference and retrieval costs can erase margins on low-priced unlimited plans.
- Integration maintenance scales faster than visible feature count.
- Community bots are inexpensive and familiar; switching requires a concrete new outcome.

- Implication: price the governed workflow and context layer, not the chat window. The underlying model should be replaceable; the permissioned operational graph should compound.

Four segments fit the long-term vision; two should define launch

Scoring uses pain urgency, Aerealith product fit, ability to reach, willingness to pay, and incumbent pressure. Five is favorable.

Segment	Core job	Pain	Fit	Reach	WTP	Incumbent gap	Priority
Technical individuals	Unify tasks, code, docs and systems	4	5	4	3	3	P1
Technical communities	Coordinate members, knowledge and projects	5	5	5	4	5	P1
Creator / game communities	Content, events, support and moderation	4	4	4	3	4	P2
Developer-led SMB teams	Context, approvals and cross-tool execution	5	4	3	5	3	P2
IT / platform teams	Knowledge plus infrastructure workflows	5	4	2	5	3	P3
Regulated enterprise	Hybrid AI with governance and audit	5	3	1	5	2	Later
Mass-market consumers	General personal companion	2	3	3	2	1	Defer

Launch segment A

Technical individual
Developer, homelab operator, maintainer, or creator who personally feels fragmentation. Low-friction PLG entry; proves personal value and source control.

Launch segment B

Technical community
Guild or project with shared knowledge and recurring coordination. Community plan creates workspace revenue and organic member distribution.

Expansion segment

Developer-led team
Five to fifty people with budget and governance needs. Requires audit, admin, onboarding, support, and stronger reliability before broad pursuit.

The best first customer already behaves like an operations hub

Aerealith should qualify accounts by cross-system complexity and recurring coordination cost, not by industry label alone.

Positive signals

High-fit account

- Discord is a daily operating surface, not only a social channel.
- GitHub, docs, calendars, tickets, and infrastructure all matter.
- The same questions recur and answers depend on changing context.
- One or more admins already maintain bots or automation.
- Mistakes, stale answers, or missed handoffs carry visible cost.
- A trusted operator can champion the product and define permissions.

Negative signals

Low-fit account

- Wants a generic chatbot with no proprietary context.
- Has a single tool and few recurring workflows.
- Expects unlimited model usage at a very low price.
- Has no owner for permissions, integrations, or escalation.
- Needs high-risk write automation before trust controls exist.
- Is primarily attracted by novelty or speculative sentience claims.

Qualification question	0 points	1 point	2 points
Active systems	1	2-3	4+
Recurring coordination workflows	None	Monthly	Weekly/daily
Shared knowledge dependency	Low	Moderate	Mission critical
Current bot / automation ownership	None	Ad hoc	Named owner
Consequence of a missed handoff	Minor	Time loss	Incident / revenue / trust
Permission complexity	Single user	Roles	Multiple roles + sensitive sources

Qualification threshold: 8+ of 12 points suggests a design-partner candidate. Favor accounts willing to share baseline time/cost and review action logs during pilots.

Three launch personas, three different definitions of value

The same product should present a different first-run proof for an individual, a community operator, and a developer-led team.

THE TECHNICAL OPERATOR

Developer / homelab / maintainer

'Show me what changed, where the answer came from, and let me approve the fix.'

WHY NOW

AI use is already habitual, but trust in accuracy is low. They value provenance and replaceable models.

ACTIVATION MOMENT

First proof: connect GitHub + one infra source; receive a cited status brief; create an approval-gated follow-up.

$8-$18 / MONTH

THE COMMUNITY LEAD

Guild owner / project manager

'Reduce repeated questions and missed handoffs without giving a bot unchecked authority.'

WHY NOW

Community scale creates repeated support and coordination work. Existing bots are narrow and fragmented.

ACTIVATION MOMENT

First proof: index permitted Discord channels + project docs; answer a recurring question; post an approved summary.

$49-$99 / WORKSPACE

THE DEVELOPER-FOUNDER

5-50 person technical team

'Give the team one governed context layer without replacing every tool we already use.'

WHY NOW

The team has budget, but suite lock-in and fragmented permissions slow deployment.

ACTIVATION MOMENT

First proof: permission-aware briefing across GitHub, docs, calendar and Discord; assign an action with an auditable receipt.

$24+ / USER

The product is hired to reduce coordination risk - not to produce more chat

These jobs translate broad vision into observable outcomes that can be tested in onboarding, pilots, and retention analysis.

When...	Help me...	So I can...	Proof / KPI
I need an answer across systems	retrieve only permitted, current context with citations	act without searching five tools or exposing the wrong data	Time-to-answer; citation open rate
A recurring workflow begins	assemble context, preview the plan, and ask for approval	delegate routine work without losing accountability	Approved action rate; exceptions
A community member asks for help	resolve from trusted project knowledge or escalate cleanly	reduce repeated support while protecting private channels	Deflection; escalation quality
Code or infrastructure changes	summarize impact and identify the owner or next action	catch drift before it becomes an incident	Detection lead time; acknowledged alerts
I revisit a long-running project	restore the relevant memory and show why it was retrieved	continue without re-explaining everything	Return-session success; memory edits
An automated action fails	show the input, decision, permissions, result, and recovery path	diagnose and recover rather than distrust the whole system	MTTR; rollback rate

Value metric

Hours of coordination avoided
Measure recurring questions, manual status gathering, handoff preparation, and incident follow-up before and after deployment.

Trust metric

Share of actions safely delegated
Track preview-to-approval rate, exceptions, permission failures, rollbacks, and unreviewed high-risk actions - which should remain zero.

Compounding metric

Context reuse with provenance
Track cited answers, repeat-source retrieval, memory corrections, source freshness, and cross-workspace context isolation.

The champion feels coordination pain; the blocker fears invisible authority

Aerealith's sales and onboarding materials should answer five stakeholder questions before an enterprise buyer asks them.

CHAMPION
Community lead / developer

Question
Will this save me time this week?

Required proof
Cited brief + first approved action

ECONOMIC BUYER
Founder / ops lead

Question
Is the saved coordination worth the recurring cost?

Required proof
Baseline hours, pilot ROI, usage guardrails

ADMIN
Platform owner

Question
Can I control sources, roles, models and integrations?

Required proof
Admin map, permission simulator, audit log

SECURITY / PRIVACY
Reviewer

Question
Where does data go, how long is it kept, and who can act?

Required proof
Data flow, retention, subprocessors, controls

END USER
Member / teammate

Question
Can I understand and correct what it knows?

Required proof
Citations, memory receipts, feedback, appeal

- GTM implication: community and team plans need a shareable trust packet before formal enterprise sales - data map, security overview, permissions model, audit examples, and model-provider policy.

The battlefield is crowded, but no category owns all four layers

Logos indicate representative products, not endorsements. Aerealith's target state sits at the intersection of context, action, community, and control.

HORIZONTAL ASSISTANTS

- ChatGPT
- Claude
- Gemini
- Copilot

KNOWLEDGE + CONTEXT

- Notion
- Atlassian Rovo
- Glean
- Guru

AUTOMATION + BUILD

- Zapier
- n8n
- Make
- Dify

COMMUNITY + DEV OPS

- Discord
- MEE6
- GitHub
- Cloudflare

- Strategic read: broad assistants win model experience; suite products win embedded distribution; knowledge vendors win enterprise permissions; automation vendors win integration breadth; community bots win familiarity. Aerealith must win the connected operating experience.

Horizontal copilots set the experience bar and compress generic-chat pricing

These products have superior distribution, model investment, and familiarity. Aerealith should interoperate with model providers while differentiating at the context, permission, community, and action layers.

C

ChatGPT Business
$20/user/mo annual; $25 monthly

Offer Company knowledge, connectors, workspace controls, strong general assistant.

Strength Deep model/product cadence and mindshare.

Aerealith opening Not community-native; product value centers on the ChatGPT workspace.

Claude Team / Enterprise
$20-$25 Team; premium seat $100-$125; Enterprise usage-based

Offer Strong writing/reasoning, projects, team and enterprise controls.

Strength Quality perception and developer affinity.

Aerealith opening Fewer community/operations primitives; actions depend on integrations and ecosystem.

M

Microsoft 365 Copilot
$18 promo / $21 regular annual; requires Microsoft context

Offer Embedded in Microsoft 365 with organizational data and admin infrastructure.

Strength Distribution and workflow proximity inside the suite.

Aerealith opening Less compelling for Discord/GitHub-first or mixed-stack communities.

Gemini in Workspace
Workspace tiers from $7/$14/$22; AI bundled by tier

Offer Gmail, Docs, Meet and Workspace-native assistance.

Strength Bundle economics and Google ecosystem reach.

Aerealith opening Cross-community and infrastructure operations are not the center of gravity.

- Response: never claim to be a better foundation model. Offer model choice and BYOK/BYOM where feasible, then make Aerealith the durable operating layer above replaceable models.

Knowledge platforms are the closest strategic threat to Aerealith's trust thesis

They already market permission-aware retrieval, citations, connectors, organizational context, and agentic workflows to enterprises.

Vendor	Public commercial model	Documented center of gravity	Strategic threat	Aerealith response
Glean	Per-user license + pooled credits; quote required	Enterprise search, permissions, cited answers, 100+ apps, agents	Very high for enterprise context	Win earlier in community/dev-led accounts; make local/hybrid and action receipts visible
Atlassian Rovo	Included in paid Atlassian Cloud; Rovo Dev $20/dev + credits	Search, chat, agents, Teamwork Graph, Jira/Confluence distribution	High in Atlassian shops	Own Discord/GitHub/infra seams; integrate rather than require suite replacement
Notion AI	AI bundled in Business; custom agents use extra credits	Workspace content, search, meeting notes, agents	High for Notion-centric SMBs	Cross-system operations, permission receipts, community workspace
Guru	Usage-shaped pricing; quote path	Governed knowledge, verification, citations, integrations	Moderate-high in support/enablement	Developer/community workflows and local runtime differentiation

Table stakes

Permission inheritance, citations, freshness, connector health, source-level controls, deletion, audit, SSO for business tiers, and defensible data-use promises.

Differentiation

A user-visible memory model, community context boundaries, GitHub/infra awareness, action preview/receipts, and a portable cloud-local-hybrid control plane.

Do not compete on

Connector count alone. Each connector adds permissions, schema drift, support burden, and security surface. Win with a smaller set of deeply instrumented workflows.

Automation platforms own breadth; Aerealith can own governed delegation

Aerealith should not recreate a visual integration builder in v1. It should provide opinionated workflows, approval semantics, and clear receipts - with escape hatches for builders.

Platform	Entry pricing	Economic unit	Strength	Gap Aerealith can exploit
Zapier	Free 100 tasks; Pro from $19.99; Team from $69	Task	9,000+ apps, mature PLG, broad templates	Builder complexity; permission/context trust is not the primary experience
n8n	Starter €20; Pro €50; Business self-host €667	Workflow execution	Developer control, self-hosting, extensibility	Requires workflow design and operations expertise
Make	Free 1,000 credits; Core $9; Pro $16; Teams $29	Credit/module action	Visual composition, 3,000+ apps, accessible price	Credit opacity and builder-first mental model
Dify	Hosted Professional $590/workspace/year; Community self-host	Workspace/resources	LLM app building, knowledge, workflows, private deployment	Platform for builders rather than finished community operations product

RECOMMENDED AUTOMATION LAYERING

NATIVE

High-frequency Aerealith workflows

Briefs, Q&A;, tasks, approvals, community posts, GitHub summaries

TEMPLATES

Curated connectors and recipes

Opinionated workflows with tested permissions and recovery

ESCAPE HATCH

Webhooks, API, MCP, n8n/Zapier handoff

Let builders extend without Aerealith owning every connector

Community bots anchor price expectations below enterprise AI

Discord-native buyers already understand bot subscriptions. Aerealith must show a step-change in coordination and context value, then price by workspace plus controlled usage.

Product	Primary job	Public price signal	AI / context model
MEE6	Moderation, levels, tickets, engagement	Promo €5.99; then €11.99 monthly; bundles vary	AI and character products may be separate subscriptions
Dyno	Moderation, automod, roles, custom commands	Free + premium structure	Modular server-management orientation
Carl-bot	Reaction roles, automod, logs, commands	Free + premium structure	Rules/command orientation
Ticket Tool	Ticketed community support	$8/month or $66/year	Narrow support workflow

200M+

DISCORD MONTHLY ACTIVE USERS

Company-reported; useful as platform scale, not addressable users.

90M+

DAILY ACTIVE USERS

Q4 2025 company figure.

Aerealith's category jump

From: a bot that performs server functions.

To: a permission-aware operating companion that understands the community's projects, code, documents, calendar, decisions, and systems - and can take approved actions across them.

The proof must be visible within one server before asking a buyer to accept a higher subscription.

Platform economics

Discord Premium Apps support user and guild subscriptions. The published growth tier is 15% of revenue for the first $1M, then the standard platform share is 30%, with processing and other transaction costs also relevant.

Implication: native app pricing needs channel-specific margin guardrails.

Personal memory creates attachment; local control creates trust - both are hard to operate

The adjacent market demonstrates demand for continuity and privacy, while also warning against dependence on a single device, model, or business model.

Product / pattern	What it proves	Risk / limitation	Aerealith lesson
Mem	Persistent personal notes, meetings and ideas can be the product	Personal scope; depends on high-quality capture and retrieval	Make memory visible, editable, attributable, and interoperable
Personal AI	Structured identity and private personal models have a dedicated audience	Claims can outrun practical daily jobs	Separate aspirational companion research from commercial proof
Limitless	Always-on memory hardware created a distinctive capture experience	Hardware/acquisition dependence; support continuity matters	Avoid a single capture modality or strategic dependency
n8n Community	Self-hosting is a credible developer distribution wedge	Operational burden shifts to the customer	Offer supported local runtime with clear responsibility boundaries
Dify Community / Enterprise	Open core can feed enterprise private deployment	Community-to-enterprise conversion requires governance and support	Open the runtime and SDK; monetize management, controls, support and scale

Recommended deployment sequence

1. Managed cloud beta with strict source/action controls.
2. Local bridge for private data and integrations.
3. Supported self-host for technical teams.
4. Hybrid / air-gapped enterprise after update, telemetry, support, and security operations mature.

Memory product contract

Every durable memory should expose source, scope, timestamp, confidence, retention, and correction. Shared workspaces need explicit ownership and deletion semantics. Memory that cannot be inspected will eventually be treated as surveillance or hallucination.

Capability scorecard: assistants and knowledge platforms

0 = absent from public materials; 5 = category-leading documented capability. Scores are directional. Aerealith is its planned target state, not a shipped-product score.

Product	Assistant	Permissioned knowledge	Visible memory	Actions	Community native	Infra aware	Hybrid / self-host	Governance
Aerealith target	4	5	5	4	5	4	5	5
ChatGPT Business	5	4	3	3	1	1	1	4
Claude Team / Ent.	5	3	3	3	1	1	1	4
M365 Copilot	4	5	2	4	1	2	1	5
Gemini Workspace	4	5	2	4	1	1	1	5
Notion Business	4	5	3	4	1	1	1	4
Atlassian Rovo	4	5	2	5	1	3	1	5
Glean	4	5	2	4	1	1	2	5
Guru	3	5	2	3	1	1	2	5

WHAT THE MATRIX SAYS

No open field

Every major capability has a strong incumbent. Aerealith's advantage can only emerge from the combination and the experience of control.

Target-state burden

The Aerealith target scores imply a large engineering and assurance surface. Sequence proof points; do not market the final state as launch scope.

Enterprise bar

Permission-aware retrieval, SSO, audit, data-use terms, and connector isolation are already table stakes for larger buyers.

Capability scorecard: automation, community, and builder platforms

Automation breadth and community familiarity are separate strengths. Aerealith should connect them through an opinionated, governed operating experience.

Product	Assistant	Permissioned knowledge	Visible memory	Actions	Community native	Infra aware	Hybrid / self-host	Governance
Aerealith target	4	5	5	4	5	4	5	5
Zapier	3	2	1	5	1	2	1	4
n8n	3	2	1	5	1	4	5	4
Make	2	1	0	5	1	2	1	3
Dify	4	3	3	5	1	2	5	4
MEE6	2	1	1	3	5	0	0	2
Dyno / Carl-bot	0	1	0	3	5	0	0	2
Ticket Tool	1	2	1	3	5	0	0	2
Mem / Personal AI	4	1	5	1	0	0	2	2

DESIGN CONCLUSION

The core product is an authority-aware context-to-action loop.

Ingest permitted context -> answer with citations -> expose memory/source scope -> preview the proposed action -> require the right approval -> execute through a bounded integration -> record the result and recovery path. This loop should feel native in Discord and the web dashboard while remaining model- and integration-agnostic underneath.

COMPETITIVE PRICING

Aerealith spans three pricing norms - bot, prosumer AI, and team software

The public ladder must tell buyers which norm applies. A single per-user price cannot efficiently cover individual, guild, team, and hybrid enterprise value.

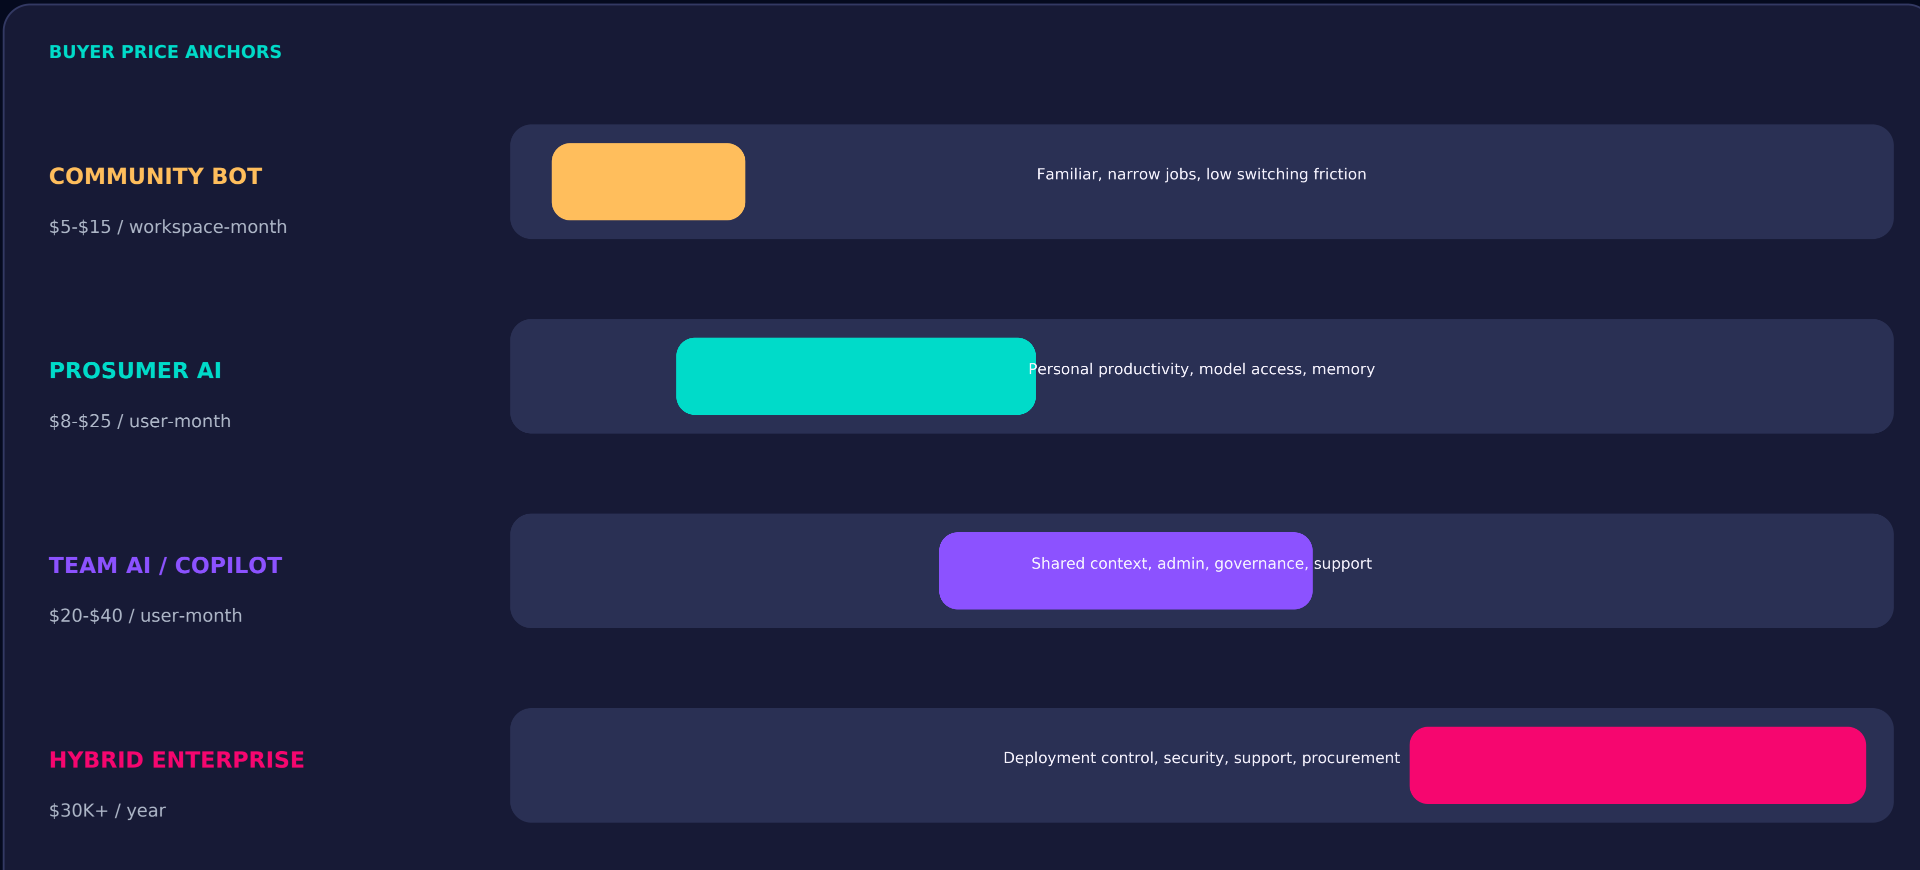

Visual bands communicate pricing archetypes; they are not plotted on a common numeric axis because enterprise ACV and seat/workspace subscriptions use different units.

- Pricing implication: preserve a low-friction personal plan, price shared community value per workspace, price team governance per seat with a minimum, and quote hybrid deployments separately.

The open position is community-native operations with enterprise-grade control

This map is an analytical synthesis of public positioning and documented capabilities, not a customer-perception survey.

Desired perception

'The AI operating layer our community can actually trust with context and approved actions.'

Avoided perception

'A Discord bot with a chatbot attached' or 'an inexpensive alternative to ChatGPT.'

Category phrase

Permission-aware AI operations
Specific enough to imply control; broad enough to expand from communities to teams.

Five compounding advantages - and five things that are not a moat

Defensibility comes from workflow adoption, permissioned context, operational trust, and ecosystem leverage. A broad feature list alone will be copied or bundled.

POTENTIAL MOATS

- Permissioned context graph - source, role, workspace, freshness and provenance relationships accumulated through real use.
- Action ledger + policy layer - reusable approval rules, receipts, recovery patterns and risk classifications.
- Community-to-organization workflow - native Discord experience connected to code, docs, calendars and systems.
- Portable deployment - consistent control semantics across managed cloud, local bridge, self-host and hybrid.
- Extension ecosystem - SDK, templates and partners that inherit Aerealith permissions and observability.

NOT MOATS

- An LLM wrapper - model experience is supplied by vendors with far greater capital and distribution.
- Low price - invites heavy usage, compresses support capacity and is easy for bundled products to undercut.
- Integration count - breadth without depth creates fragile permissions and maintenance burden.
- Generic memory - a vector database is infrastructure, not user trust or differentiated continuity.
- Sentience / NHP language - difficult to substantiate and risky as the primary commercial promise.

- Moat test: does the advantage improve as more trusted workflows run, remain valuable when the model changes, and become expensive for a customer to recreate without locking up their data?

A bottom-up market boundary is more useful than a trillion-dollar AI headline

The model counts plausible paying accounts and multiplies them by a target annual contract value. It deliberately excludes hyperscaler infrastructure, model training, advertising, hardware, and unrelated AI applications.

$64B
AI MODELS + PLATFORMS, 2026
Gartner category climate; not Aerealith TAM.

180M+
GITHUB DEVELOPERS
Audience proxy; not paying users.

36.2M
U.S. SMALL BUSINESSES
Business-universe proxy; not addressable accounts.

1. DEFINE BUYER

Technical individual, community workspace, developer-led team, or hybrid enterprise.

2. SET ACCOUNT BASE

A modeled subset of broad platform and business populations that fits the job.

3. APPLY ARR

Realistic annual pricing for the unit: user, workspace, team, or contract.

4. STRESS TEST

Use conservative, base, and upside capture scenarios; validate with pilots.

MODEL RULES

- TAM is the targetable long-run category boundary under Aerealith's strategy - not every person who uses Discord, GitHub, or AI.
- SAM applies a five-year geographic, product, channel, and delivery constraint.
- SOM is expressed as account-count scenarios and ARR, making every assumption editable.
- Account groups are designed as different commercial units, but real-world overlap remains; a sensitivity haircut should be used in fundraising models.

Modeled targetable market: $3.05B annual recurring revenue

This strategic TAM assumes Aerealith eventually serves four distinct commercial units. It is a long-run planning boundary, not a forecast or third-party market estimate.

Commercial unit	Modeled addressable accounts	Annual revenue unit	Formula	Targetable ARR
Paid technical individuals	10.0M users	$120 / user	10.0M x $120	$1.20B
Technical community workspaces	250K workspaces	$600 / workspace	250K x $600	$150M
Developer-led small teams	500K teams x 10 seats	$240 / seat	500K x 10 x $240	$1.20B
Hybrid / enterprise	10K organizations	$50K / contract	10K x $50K	$500M
TOTAL TARGETABLE TAM	-	-	Sum	$3.05B

TAM COMPOSITION

Revenue shares: individuals 39% | communities 5% | teams 39% | enterprise 16%

Critical caveat: GitHub, Discord, and SBA counts establish that large reachable populations exist; they do not validate the modeled account counts. Validate willingness to pay and segment overlap before using this TAM in external materials. A 20% overlap/eligibility haircut would reduce the boundary to approximately $2.44B.

Five-year serviceable market: $460M; base capture scenario: $6.98M ARR

The SAM narrows geography, product readiness, and distribution. The SOM scenarios translate strategy into account goals that can be tested quarter by quarter.

SAM unit	Accounts	ARR unit	SAM ARR
Individuals	1.50M	$120	$180M
Communities	50K	$600	$30M
Teams	75K	$2,400	$180M
Enterprise	2K	$35K	$70M
TOTAL	-	-	$460M

5-year scenario	Accounts mix	ARR	SAM capture
Conservative	6K individuals; 400 communities; 400 teams; 10 enterprise	$2.27M	0.49%
Base	15K individuals; 1.5K communities; 1.2K teams; 40 enterprise	$6.98M	1.52%
Upside	30K individuals; 4K communities; 2.5K teams; 100 enterprise	$15.50M	3.37%

Interpretation: the base outcome does not require mass-market adoption. It requires strong retention in a narrow technical wedge, repeatable community distribution, and a modest number of higher-value teams and hybrid contracts.

The base scenario depends more on team expansion than consumer scale

A mixed model lowers dependence on any one channel. Individual and community products create reach; team and enterprise contracts create revenue density.

ARR SCENARIOS BY REVENUE STREAM ($M)

- Sensitivity: if team ACV is 25% lower and individual churn prevents scale beyond 10,000 paid users, the base mix falls to roughly $5.2M ARR unless community or enterprise revenue compensates.

The internal pricing ladders are thoughtful cost controls - but too complex for launch

Aerealith currently documents two different public ladders. Both under-communicate the distinct value of a personal account, a community workspace, and a governed team deployment.

Internal model	Tiers	Price points	Strength	Risk
Pricing Model v3	8 including Enterprise	$0 / $3 / $6 / $10 / $15 / $20 / $25 / custom	Granular usage segmentation; explicit credits/automations/API/storage	Choice overload, thin differentiation, $3 increments amplify support and COGS risk
MVP Master Plan	5 including Enterprise	$0 / $5 / $10 / $20 / waitlist	Simple, recognizable personal SaaS ladder	No community workspace unit; unclear team/enterprise bridge

Problem 1: unit

The price changes by personal feature/usage tier, while shared community and team value are not expressed as separate buying units.

Problem 2: margin

Low fixed prices plus AI credits, storage, API requests, automation, support and connector maintenance can produce unbounded cost-to-serve.

Problem 3: message

Eight tiers make prospects compare quotas before they understand the core outcome: trusted context and approved action across systems.

Pricing principle: package by the unit of shared value, then meter the expensive resources underneath. Buyers should choose based on who uses Aerealith and what governance they need; model tokens, large retrieval workloads, storage, and high-volume automations should be governed by credits, fair-use limits, BYOK, or overage.

Recommended launch packaging: four public offers, one enterprise path

The numbers are test prices. Validate through design-partner willingness-to-pay interviews and live conversion; preserve annual discounts and explicit usage guardrails.

FREE

$0
Personal

1 workspace; small monthly AI/memory cap; 1 core integration; cited answers; manual approvals

Role: Activation + trust proof

PERSONAL

$8/mo
Individual

More context and memory; 3-5 integrations; tasks; summaries; BYOK option

Role: Habit + continuity

POWER

$18/mo
Technical individual

GitHub + Discord + read-only infra; workflows; advanced memory; API/webhooks

Role: Technical wedge

COMMUNITY

$49/mo
Guild / project workspace

Shared permitted knowledge; roles; community briefs; approved posts/actions; admin usage pool

Role: Workspace distribution

TEAM

$24/user/mo
5+ person organization

Shared context; admin; policy templates; audit; SSO roadmap; support; pooled usage

Role: Expansion revenue

ENTERPRISE HYBRID

$30K-$50K+ ARR
Private deployment

Hybrid/self-host, SSO/SCIM, audit export, SLA, security review, support and implementation

Role: Contract density

- Optional Community Growth tier: $99/month with a larger pooled allowance, multiple admins, higher-frequency scheduled briefs, and support. Avoid adding it until $49 customers consistently hit a real value threshold.

Usage controls are product design, not a hidden finance mechanism

Aerealith needs transparent limits, model routing, caching, retrieval discipline, and channel-aware pricing to preserve trust and gross margin.

>75%
BLENDED CLOUD GROSS MARGIN
Management guardrail, not a public promise.

≤20%
AI COGS / REVENUE
Management guardrail, not a public promise.

≤25%
TOTAL VARIABLE COGS
Management guardrail, not a public promise.

<$0.50
FREE MAU MONTHLY COGS
Management guardrail, not a public promise.

Cost driver	Control	User-facing expression	Escalation
Model inference	Route by task; cache; prompt budgets; BYOK/BYOM	Monthly AI allowance + quality modes	Paid credit pack or own provider key
Retrieval / indexing	Source limits; incremental sync; freshness windows	Connected-source and indexed-content allowance	Workspace add-on
Automation execution	Risk class; concurrency; retry caps; idempotency	Included approved actions / workflow runs	Usage pack; enterprise pool
Storage + memory	Retention tiers; compression; user deletion	Clear storage and retention controls	Storage add-on
Support + integration maintenance	Supported connector matrix; diagnostics; SLA tiers	Community vs priority support	Implementation / support plan
Discord distribution	Model 15% growth-tier platform share + processing	Channel-specific package math	Direct team/enterprise motion

North-star economic loop: more trusted workflows -> higher retention and workspace expansion -> richer permissioned context -> better answers and templates -> more trusted workflows. If usage grows without retention, expansion, or proven time savings, the plan is subsidizing activity rather than creating value.

Profit is a sequence of gates, not one forecast number

The model moves from price to cost-to-serve, contribution, customer acquisition, fixed operating expense, break-even, EBITDA and cash need. All dollar values are USD and pre-tax.

$6.98M
BASE YEAR-5 EXIT ARR
Scenario output, not guidance.

77%
BASE YEAR-5 GROSS MARGIN
Scenario output, not guidance.

$1.05M
BASE YEAR-5 EBITDA
Scenario output, not guidance.

$2.9M
MODELED FUNDING ENVELOPE
Scenario output, not guidance.

01 REVENUE
Price x active paid units; enterprise recognized ratably.

02 GROSS PROFIT
Revenue less inference, hosting, payment, support and direct delivery.

03 CONTRIBUTION
Gross profit less variable acquisition and channel-specific cost.

04 EBITDA
Contribution less R&D;, sales, G&A; and other fixed operating expense.

05 CASH
EBITDA adjusted for timing, capex, working capital and safety buffer.

Decision rule: do not scale a plan because revenue rises. Scale only when (1) gross margin remains above the plan floor, (2) cohort payback is acceptable, (3) support and connector labor is measured, and (4) retention proves the workflow is durable. A fast-growing low-margin tier can increase cash burn faster than it increases enterprise value.

Accounting boundary: modeled EBITDA is not net income. It excludes interest, taxes, depreciation, amortization, stock compensation, fundraising costs and jurisdiction-specific obligations. The funding envelope is a planning estimate, not a financing recommendation.

Cost-to-serve: instrument six meters before public launch

Provider prices are snapshots; Aerealith's actual cost depends on token mix, cache rate, retrieval design, concurrency, support burden and customer behavior.

Meter	Cost basis	Base planning method	Control / owner
Inference	Input + output + cached tokens by model	Per workflow and account; premium reasoning separated	Router, prompt budget, cache / AI Eng
Retrieval + indexing	Embeddings, sync, database, storage, egress	Incremental sync and source-level allocation	Freshness window, caps / Data
Automation	Requests, CPU, retries, external API calls	Successful + failed execution by workflow	Idempotency, retry cap / Platform
Payments + channel	Processor, refund, dispute, tax and platform share	Direct and Discord P&L; kept separate	Annual billing, channel price / Finance
Support	Loaded labor minutes per account and connector	Ticket + implementation time attributed by plan	Docs, diagnostics, SLA / Success
Free usage	All direct delivery cost for non-paying MAUs	Free COGS per MAU and per converted payer	Hard monthly cap / Growth

Model pricing is nonlinear

Different models, long context, search tools, caching, batch and priority modes produce materially different unit costs. Route on task value, not brand preference.

Infrastructure has steps

Low-volume cloud allowances can make beta cost look artificially cheap; enterprise controls and observability create later fixed-cost steps.

Labor is real COGS

Implementation, connector troubleshooting and premium support must be attributed to the customer or the reported software margin will be overstated.

- Minimum telemetry: cost per request, active account, successful workflow, cited answer, approved action and retained account - segmented by plan, model, connector and acquisition channel.

Plan-level cost-to-profit: viable at launch, but below mature SaaS margin

Modeled monthly economics use recommended prices, representative usage, direct US card processing at 2.9% + $0.30, and attributed support. Team assumes eight seats; enterprise assumes $35K ARR.

Plan / unit	Revenue	AI	Infra + data	Support	Payment	Total COGS	Gross profit	GM
Personal account	$8.00	$0.90	$0.80	$0.40	$0.53	$2.63	$5.37	67%
Power account	$18.00	$2.60	$1.85	$0.90	$0.82	$6.17	$11.83	66%
Community direct	$49.00	$5.50	$4.50	$3.50	$1.72	$15.22	$33.78	69%
Team workspace / 8 seats	$192.00	$18.00	$16.00	$10.00	$5.87	$49.87	$142.13	74%
Enterprise recurring / month	$2,917	$220	$340	$270	$30	$860	$2,057	71%

GROSS-PROFIT DOLLARS PER MONTH

Interpretation: launch margins of 66-74% are plausible for an AI-native product but trail Benchmarkit's 81% median subscription gross margin. Aerealith should target 75% blended margin by Year 4 through routing, caching, annual billing, support automation and mix shift - not by hiding hard usage caps.

- Enterprise onboarding is excluded from recurring COGS: charge a separate implementation fee and measure that service margin independently. Free-plan cost is analyzed on the next page.

Free and Discord economics can quietly erase contribution

The same $49 Community product has different profit depending on channel. Free usage is a portfolio acquisition cost, not zero-cost revenue.

$49 Community month	Direct web	Discord growth tier	Discord standard tier
Gross revenue	$49.00	$49.00	$49.00
Payment / platform cost	$1.72	$7.35 + processing*	$14.70 + processing*
AI + infra + support	$13.50	$13.50	$13.50
Modeled contribution	$33.78	$26.68 before processing	$19.33 before processing
Contribution margin	69%	54% before processing	39% before processing

FREE-TO-PAID SUBSIDY

At $0.35 per active free user/month:

Free MAUs per payer	Subsidy / payer	Personal GM after subsidy
2	$0.70	58%
3	$1.05	54%
5	$1.75	45%
10	$3.50	23%

Channel rule

Price native subscriptions separately or constrain the included usage pool. Do not assume direct-web unit economics survive platform share.

Free rule

Keep free COGS below $0.50/MAU and graduate inactive accounts to on-demand reactivation. Measure conversion by retained payer, not signup.

Profit protection sequence
1. Show users their allowance and expensive actions. 2. Route summaries and classification to lower-cost models. 3. Cache stable context and sync incrementally. 4. Bill annually where fit. 5. Apply BYOK to exceptional workloads. 6. Pause or reprice workflows whose usage grows without retention or expansion.

*Discord documentation describes platform share; transaction, tax, payout, refund and regional mechanics require launch-date confirmation. Processing is deliberately not guessed here.

CAC, LTV and payback: acquisition ceilings by plan

Simple steady-state model: gross-profit LTV = monthly revenue x gross margin / monthly logo churn. It excludes expansion, discounting and survival-curve effects and should be replaced by cohort data.

Plan	Monthly GP	Churn assumption	Gross-profit LTV	CAC ceiling used	LTV:CAC	Payback
Personal	$5.37	5.0% / mo	$107	$32	3.4x	6.0 mo
Power	$11.83	4.0% / mo	$296	$90	3.3x	7.6 mo
Community	$33.78	3.0% / mo	$1,126	$300	3.8x	8.9 mo
Team / 8 seats	$142.13	1.5% / mo	$9,475	$1,200	7.9x	8.4 mo
Enterprise	$24.7K / yr	15% / yr	$164K	$25K	6.6x	12.2 mo

PLG ceiling

Keep Personal and Power blended payback under 9 months. Paid media is unlikely to work if it needs sales assistance or support-heavy onboarding.

Workspace ceiling

Community and Team can support higher CAC only when member invitations, templates and retained shared workflows create expansion.

Enterprise ceiling

A $25K CAC requires implementation fees, strong qualification and a repeatable security package. One-off custom engineering belongs in deal economics.

Required cohort dashboard: CAC by channel and segment; activation; week-4 and month-3 retention; gross revenue retention; net revenue retention; gross margin; expansion; support minutes; payback; and LTV based on observed survival. Benchmarkit's margin data and SaaS Capital's 91% GRR / 103% NRR are context, not substitutes for Aerealith cohorts.

Break-even depends more on mix than raw customer count

Break-even paid units = monthly fixed operating expense / weighted monthly contribution. A unit is one paid account or workspace; enterprise contracts are shown separately.

Operating stage	Monthly fixed OPEX	At $12 contribution	At $18 contribution	At $30 contribution
Founder / lean	$35K	2,917 units	1,945 units	1,167 units
Base team	$85K	7,084 units	4,723 units	2,834 units
Growth team	$180K	15,000 units	10,000 units	6,000 units

BASE-TEAM BREAK-EVEN COMBINATIONS

Illustrative mix	Monthly contribution	Remaining OPEX gap	Additional need
3,000 individuals @ $8 avg GP	$24,000	$61,000	1,806 Community direct
1,000 Community direct	$33,780	$51,220	361 Team workspaces
300 Team workspaces	$42,639	$42,361	21 Enterprise contracts*
20 Enterprise contracts*	$41,140	$43,860	1,299 Community direct
Balanced: 2K indiv + 700 Community + 200 Team + 8 Ent.*	$84,161	$839	~47 more avg units

Management conclusion: the shortest path to break-even is not maximizing $8 subscribers. It is earning a balanced revenue mix in which Community validates workspace distribution, Team adds governance economics, and a small number of qualified enterprise contracts contributes without consuming custom-engineering capacity.

*Enterprise contribution is monthly recurring gross profit per active annual contract. Sales capacity, implementation timing and collections still affect cash break-even.

FINANCIAL MODEL

Base five-year P&L; reaches EBITDA break-even in Year 4

Illustrative operating model tied to the report's base SOM. Revenue is approximated from average beginning/end ARR; gross margin expands with routing, annual billing, support leverage and mix.

USD millions	Year 1	Year 2	Year 3	Year 4	Year 5
Exit ARR	$0.25	$0.90	$2.10	$4.20	$6.98
Recognized revenue	$0.13	$0.58	$1.50	$3.15	$5.59
Gross margin	65%	69%	72%	75%	77%
Gross profit	$0.08	$0.40	$1.08	$2.36	$4.30
Operating expense	($0.80)	($1.25)	($1.75)	($2.35)	($3.25)
Modeled EBITDA	($0.72)	($0.85)	($0.67)	$0.01	$1.05
EBITDA margin	-576%	-148%	-45%	0%	19%

EBITDA PATH

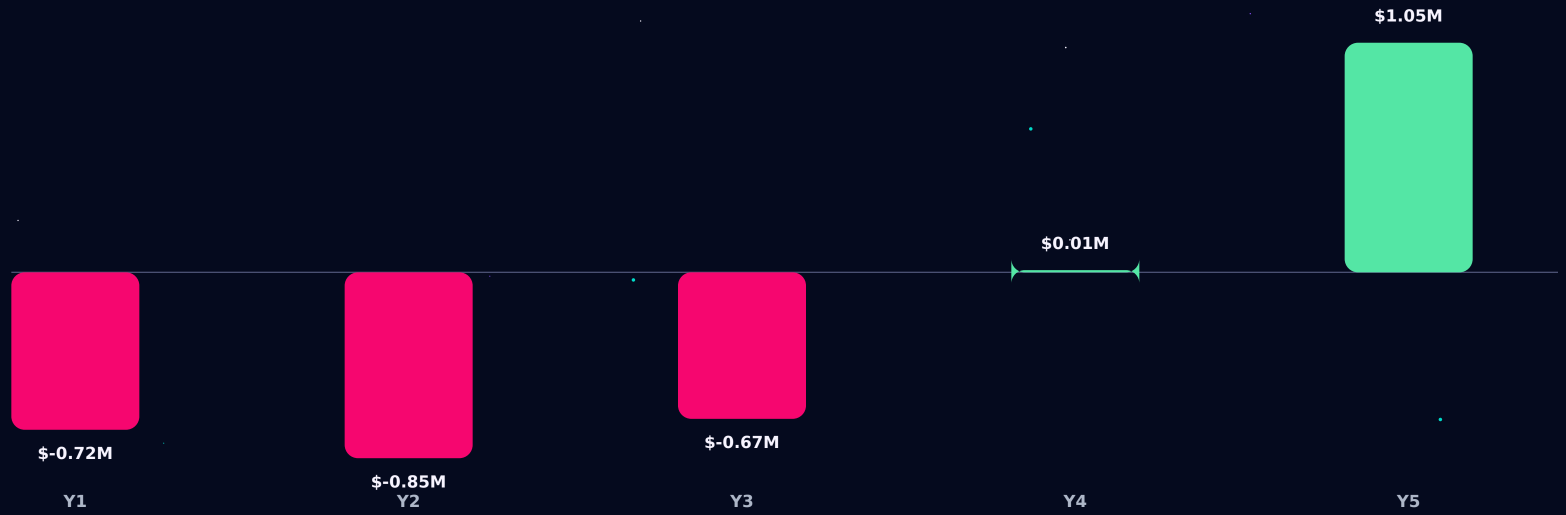

- The model is intentionally conservative on early margin and excludes tax benefits from losses. It is a capacity and funding plan, not audited guidance or a valuation model.

Year-5 profit scenarios: retention and mix separate survival from scale

Three planning cases isolate strategic outcomes. Scenario values are modeled, not probabilities. The downside case requires earlier scope and cost action.

Year-5 metric	Conservative	Base	Upside
Exit ARR	$2.27M	$6.98M	$15.50M
Recognized revenue	$1.85M	$5.59M	$12.10M
Gross margin	68%	77%	78%
Gross profit	$1.26M	$4.30M	$9.44M
Operating expense	$2.00M	$3.25M	$6.50M
EBITDA	($0.74M)	$1.05M	$2.94M
EBITDA margin	-40%	19%	24%

CONSERVATIVE

Driver: Low conversion, weak expansion, 68% margin

Action: Cut connector breadth; keep team lean; require paid implementation; seek break-even niche or pause.

BASE

Driver: Community-to-Team expansion; 77% margin

Action: Fund trust surface and repeatable GTM; maintain hard CAC and payback gates.

UPSIDE

Driver: Strong retention, enterprise pull, efficient routing

Action: Add capacity behind proven segments; prevent custom deals from eroding margin and roadmap.

- Profit should be reviewed with growth, GRR/NRR, payback and cash. A positive EBITDA number created by under-investing in reliability or support is not a durable outcome.

FINANCIAL MODEL

Cash plan: roughly $2.9M before financing and contingency choices

Base case accumulates approximately $2.24M of EBITDA losses through break-even. A 30% operating and timing reserve produces a $2.9M planning envelope.

FUNDING ENVELOPE BRIDGE

$2.90M
BASE FUNDING ENVELOPE
Modeled; excludes financing fees.

Q4 YEAR 4
BASE EBITDA BREAK-EVEN
Timing varies with hiring and collections.

30%
RESERVE ABOVE PEAK LOSS
For volatility, not planned spend.

Stage capital against evidence: (1) fund 90-day validation and trust architecture; (2) fund 12-18 months only after retained paid cohorts; (3) expand sales and enterprise controls after repeatable payback. Monthly reforecast should include cash balance, committed payroll, contracted cloud minimums, receivables timing, annual prepayments, refunds and downside runway.

- This is a management planning envelope, not a recommendation to raise a specific security or amount. Financing structure, dilution and taxes require professional advice.

Sensitivity: churn, price realization and mix dominate inference savings

Directional Year-5 EBITDA effects around the base case. One variable changes at a time; interactions can amplify the result.

Variable change vs base	Approx. EBITDA effect	Management response
Realized ARPA -10%	-$0.56M	Reduce discount leakage; repackage high-value workflows
Gross margin -5 points	-$0.28M	Route/cache; fix support-intensive connectors; price heavy usage
Recognized revenue -20%	-$0.86M GP	Slow hiring; protect trust roadmap; cut unproven channels
OPEX +15%	-$0.49M	Stage hiring to retention and payback gates
Enterprise mix +5 points	+$0.20M to +$0.45M	Only if implementation and sales capacity remain repeatable
Free COGS doubles	-$0.05M to -$0.20M	Tighten inactivity, model routing and conversion qualification

PRIORITY ORDER

1 RETENTION + ARPA

Validate durable value and willingness to pay before traffic.

2 MIX + PAYBACK

Move successful communities into team governance without forced upsell.

3 DELIVERY MARGIN

Instrument inference, support and connectors; optimize the expensive workflow.

4 FIXED COST

Hire after gates; keep reliability and security non-discretionary.

- Inference optimization matters, but it cannot rescue a product with weak retention or low price realization. The model should be refreshed monthly with actual cohorts and provider invoices.

Finance operating system: twelve controls for the next 90 days

The immediate objective is not precision theater. It is a repeatable cost and cohort ledger that makes pricing, product scope and hiring decisions observable.

01 **Tag every request**
Account, plan, workflow, model, tokens, cache, connector and outcome.

02 **Attribute infra**
Database, vector, storage, egress and automation to active accounts.

03 **Time support**
Ticket and implementation minutes by plan and connector.

04 **Separate channels**
Direct, Discord and enterprise deal P&Ls; never blend platform share away.

05 **Define COGS**
Publish the accounting policy for support, DevOps and implementation.

06 **Set plan floors**
Alert at 60% individual, 65% workspace and 70% team recurring GM.

07 **Gate free usage**
Cost/MAU, activation, retained conversion and subsidy per payer.

08 **Track cohorts**
Activation, W4/M3 retention, GRR, NRR, expansion and gross profit.

09 **Measure CAC**
Fully loaded channel CAC; exclude organic only when genuinely zero spend.

10 **Reforecast cash**
Monthly hiring, cloud commitments, annual prepay, receivables and reserve.

11 **Price exceptions**
Every override has an owner, expiry, usage cap and contribution test.

12 **Review monthly**
Product, engineering, growth and finance decide routing, packaging and spend.

- Owner recommendation: one accountable finance/operations lead maintains definitions; engineering owns metering quality; product owns plan guardrails; growth owns cohort acquisition economics.

Land through individual operators; expand through shared workflows

The recommended motion combines product-led adoption, community-native distribution, open-source credibility, and founder-led design partnerships.

DISCOVER

Technical operator sees a workflow-specific demo or installs the Discord app.

Motion:
Content, app directory, GitHub, referrals

ACTIVATE

Connect two sources; receive one cited answer; approve one bounded action within seven days.

Motion:
Guided setup + templates

ADOPT

Use weekly brief, project Q&A;, handoff, or incident workflow repeatedly.

Motion:
Recurring value loop

EXPAND

Invite members, add shared context, upgrade community or team governance.

Motion:
Workspace collaboration

STANDARDIZE

Adopt policy templates, audit, local bridge or hybrid deployment.

Motion:
Team / enterprise motion

Channel	Role	First asset	Success signal	Risk
Discord App Directory / guild installs	Reach technical communities	Permission-aware project brief	Qualified installs -> activation	Low-price bot comparison
Developer content + GitHub	Reach operators and maintainers	Open local bridge, workflow templates, changelog	Install, star, template reuse, signup	Open-source users may not convert
Founder-led design partners	Learn high-value workflow	Measured before/after pilot	4+ weekly-active pilots; paid conversion	Custom work creates roadmap drift
Community case studies	Reduce trust uncertainty	Action receipt, time saved, incident avoided	Referral installs and higher conversion	Requires credible evidence
Integration partners	Extend reach without owning every connector	Certified adapter + shared launch	Activated joint accounts	Support and permission ambiguity

- Sequence rule: do not scale acquisition until the activation loop and week-4 retention work in a narrow workflow. Distribution amplifies product truth - good or bad.

Discord is both the beachhead interface and a margin-bearing channel

Native discovery and guild subscriptions can accelerate reach. Published platform economics must be included in pricing, COGS, and packaging decisions.

Modeled $49 Community month	Growth tier	Standard tier
Gross price	$49.00	$49.00
Platform share	15% = $7.35	30% = $14.70
Illustrative processing assumption	3% = $1.47	3% = $1.47
Net receipts before Aerealith COGS	$40.18	$32.83
Illustrative variable COGS	$8.00	$8.00
Contribution before fixed cost	$32.18	$24.83
Contribution / gross revenue	65.7%	50.7%

Processing is a modeling assumption, not a Discord quote. Actual transaction, tax, refund, and regional costs vary. Confirm current platform terms before launch.

Decision: test $49 direct economics and a higher native-channel price or tighter usage pool. The current low-price internal tiers are not robust to platform share plus model and support cost.

Use native distribution for

Guild installs, trial, community-plan purchase, member invitations, recurring briefs, and workflows where Discord is the operating surface.

Use a distinct web business motion for

Team administration, complex procurement, hybrid deployment, implementation, security review, and contracts whose product scope exceeds a Discord app subscription.

Do not do

Design pricing to evade platform policy, hide fees, or make users guess where a purchase belongs. Treat platform compliance and customer clarity as product requirements.

Lead with the governed outcome; reveal the larger vision after trust is earned

Aerealith's brand is emotionally expansive. The launch message should retain that warmth while becoming concrete enough for a technical buyer to verify.

RECOMMENDED HOMEPAGE FRAME

The permission-aware AI operations layer for your community and tools.

Ask across Discord, GitHub, documents, calendars, and systems. See every source. Control every memory. Approve every consequential action.

CONNECT TWO SOURCES

SEE A LIVE ACTION RECEIPT

Message pillar 1

Context you can verify
Permission-aware retrieval, visible citations, source freshness, and connector health.

Message pillar 2

Memory you control
Scope, origin, retention, correction, deletion, and shared-workspace ownership.

Message pillar 3

Actions you authorize
Preview, risk class, approval, bounded execution, receipt, and recovery.

De-emphasize at launch: 'everything app,' generalized command center claims, unlimited integrations, autonomous operation, and NHP/personhood language. These widen the promise faster than the proof and invite comparison with products Aerealith should not fight head-on.

Build the trust primitives before the feature perimeter

The market review changes roadmap order. A small number of deeply instrumented workflows will create more differentiation than a broad connector catalog.

P0 CONTROL PLANE

- Source and workspace permission model
- Citations + freshness + connector health
- Memory receipts, editing, deletion and retention
- Action preview, approval, idempotency and ledger
- Data-use policy and model-provider boundaries

Release blocker. The product promise is false without it.

P1 BEACHHEAD OUTCOMES

- Discord + GitHub + docs context brief
- Recurring project / community Q&A;
- Approved status post and handoff creation
- Read-only Proxmox / service health summary
- Task and calendar follow-through

Activation and retention engine for launch users.

P2 PLATFORM LEVERAGE

- Webhooks, API and MCP surface
- Local bridge and BYOK/BYOM
- Template/versioning system
- Integration SDK and permission contract
- Usage telemetry, billing and admin diagnostics

Scale and ecosystem leverage after the core loop works.

- Scope cut: defer Gmail write actions, financial integrations, broad infrastructure writes, marketplace economics, and high-autonomy agents until the permission/action model survives real pilots.

Twelve months to a credible, paid beachhead

Dates begin when a dedicated product team is staffed. Advancement depends on evidence gates, not calendar completion alone.

0-90 DAYS

FOUNDATION

- 15-20 buyer interviews
- 5 design partners
- Permission/source schema
- Memory receipt prototype
- Action ledger + risk classes
- Unified launch pricing decision

Evidence gate
Gate: 4 partners activate; no critical permission defect

MONTHS 4-6

PRIVATE BETA

- Discord + GitHub + docs
- Cited project brief
- Approved post/task action
- Read-only infra summary
- Usage/cost telemetry
- Weekly pilot reviews

Evidence gate
Gate: 40%+ week-4 workspace retention

MONTHS 7-9

PAID EARLY ACCESS

- Personal / Power / Community
- Billing + pooled usage
- Templates + diagnostics
- Export/delete controls
- Trust packet + status page
- First public case studies

Evidence gate
Gate: 25+ paid; positive contribution cohort

MONTHS 10-12

REPEATABLE LAUNCH

- App-directory optimization
- Referral/invite loop
- Team admin beta
- Local bridge alpha
- Partner adapter pilot
- Reliability and support runbook

Evidence gate
Gate: repeatable activation and <6-mo PLG payback

- The roadmap assumes read-only infrastructure at launch. Write access to production systems should remain separately risk-classified and approval-gated until recovery and incident response are proven.

Expansion follows the same primitives into larger operating contexts

Aerealith should expand by reusing identity, permissions, memory, policies, and receipts - not by launching disconnected vertical products.

12-18 MONTHS
TEAM

Team workspaces, policy templates, admin analytics, SSO beta, shared memory ownership, support tiers, integration health.

18-24 MONTHS
LOCAL + HYBRID

Supported local bridge, model-provider policy, private connectors, self-host preview, deployment diagnostics, signed updates.

24-30 MONTHS
ECOSYSTEM

SDK, certified adapters, template marketplace, partner support boundaries, versioned permissions, revenue-share experiment.

30-36 MONTHS
ENTERPRISE

SCIM, audit export, regional controls, private deployment, SLA, security program, procurement package, implementation partners.

Expansion gate: do not enter regulated or large-enterprise procurement until Aerealith can produce a current data-flow map, security controls, subprocessors, model-use policy, deletion/retention evidence, audit samples, incident process, and a clear shared-responsibility model for each deployment mode.

Turn permission-aware from a claim into a visible control plane

Every answer and action should carry a machine-readable and human-readable chain of authority. This becomes the product's trust signature.

SOURCE PLANE	System: Discord channels, GitHub repos, docs, calendars, systems	Visible control: Identity, scope, freshness, connector health
CONTEXT PLANE	System: Retrieval, memory, summarization, citations	Visible control: Source receipt, confidence, retention, correction
POLICY PLANE	System: Workspace roles, risk classes, model/provider rules	Visible control: Who may see, propose, approve, execute, or export
ACTION PLANE	System: Preview, approval, bounded execution, validation	Visible control: Idempotency, result receipt, rollback / recovery
ASSURANCE PLANE	System: Audit, evaluation, incidents, user appeal	Visible control: Metrics, alerts, red-team cases, deletion evidence

- Launch acceptance rule: a user should be able to answer 'why did it know this?' and 'why was it allowed to do that?' without reading logs or contacting support.

The regulatory direction rewards visible governance and punishes unsupported claims

This is strategic product guidance, not legal advice. Obtain counsel for launch jurisdictions, contracts, data processing, and sector-specific use cases.

Framework / authority	Current signal	Aerealith implication	Build evidence
EU AI Act	Broad application began 2 Aug 2026; some high-risk dates extend into 2027-2028	Classify use cases; maintain transparency, literacy, documentation and provider/deployer responsibilities	Use-case inventory, risk classification, notices, technical documentation
NIST AI RMF + GenAI Profile	Voluntary Govern / Map / Measure / Manage structure	Use as operating system for product risk and enterprise trust packet	Risk register, evaluations, incident and change process
U.S. FTC	Enforcement targets deceptive AI claims and broken privacy/confidentiality promises	Substantiate accuracy, capability, earnings/productivity and data-use statements	Claim review, test evidence, privacy promise controls
UK ICO / GDPR principles	Personal data in AI requires lawful, fair, transparent lifecycle handling	Map data, purpose, retention, rights, profiling and agent responsibility	DPIA triggers, data flow, deletion and explanation tools

Launch minimum

AI disclosure, privacy/data-use policy, retention/deletion controls, model-provider list, permission design, action audit, incident contact, acceptable-use boundaries.

Contract minimum

DPA where applicable, subprocessor terms, security schedule, service boundaries, support and incident terms, customer configuration responsibilities.

Claims rule

Never claim perfect accuracy, autonomous safety, legal/compliance outcomes, sentience, or privacy properties that the implemented system and vendors cannot prove.

Keep NHP and companion research separate from the launch value proposition

Aerealith can explore durable identity and richer companionship, but commercial messaging should remain anchored in observable utility, user agency, and psychological safety.

Commercial product contract

DO

- Disclose that the user is interacting with AI.
- Let users inspect, correct, scope, export and delete memory.
- Make uncertainty, sources and action boundaries visible.
- Design age-aware defaults and guardian/admin controls where relevant.
- Provide escalation and crisis-safe behavior for sensitive interactions.
- Measure healthy utility, not emotional dependency.

Research boundary

DO NOT

- Market unverified sentience, consciousness or personhood claims.
- Encourage exclusivity, guilt, coercion, jealousy or dependency.
- Blur system incentives, sponsored content, or memory behavior.
- Use minors or vulnerable users as an uncontrolled experimentation cohort.
- Let a companion bypass workspace permissions or action approval.
- Treat anthropomorphic engagement as proof of wellbeing or truth.

RECOMMENDED GOVERNANCE

Create a separate research charter, review board, risk taxonomy, participant consent process, age and vulnerability protections, red-team program, publication standard, and kill criteria. NHP research may shape long-term memory and identity design; it should not be required to believe the near-term commercial case.

UTILITY FIRST | AGENCY ALWAYS | CLAIMS REQUIRE EVIDENCE

SWOT: the vision is differentiated; execution breadth is the primary internal threat

The highest-leverage strategic move is narrowing the first promise while preserving the architecture needed for later expansion.

STRENGTHS

- Distinct combination of community, developer, infra, memory and action concepts
- Strong visual identity and emotionally resonant brand
- Local-first / hybrid ambition fits trust-conscious technical users
- Safety-by-default and approval-gated direction in the internal plan
- Founder-market proximity to game, creator and technical communities

WEAKNESSES

- MVP scope spans many difficult product and security surfaces
- Pricing, hosting and audience statements conflict across internal documents
- Limited incumbent distribution, trust evidence and customer proof
- Connector and operations burden can exceed a small team's capacity
- NHP language can distract from a practical launch case

OPPORTUNITIES

- High AI adoption with weak scaling and trust maturity
- Discord/GitHub-first communities are underserved by suite copilots
- Permission receipts and action ledgers can become a recognizable product signature
- Open local runtime can seed developer adoption and hybrid demand
- Community workspace can bridge PLG to team revenue

THREATS

- Microsoft, Google, OpenAI and Atlassian bundle agents into existing spend
- Knowledge vendors already lead permission-aware enterprise retrieval
- Community bots anchor low price and have existing install bases
- Inference, support and platform fees pressure low tiers
- Privacy, security, parasocial or action incidents could damage a young brand

- SWOT action: use strengths to win a technical-community niche; use the niche to generate proof; use proof to fund control-plane maturity; expand only when the same primitives support the next segment.

Five forces: attractive demand, difficult economics, high substitution

The category is investable only with a narrow wedge, disciplined COGS, and a control layer that remains valuable across models and integrations.

VERY HIGH — **RIVALRY**

Suite incumbents, model vendors, knowledge platforms, automation builders, bots and open-source tools overlap.

VERY HIGH — **SUBSTITUTES**

Manual search, generic chat, native suite AI, scripts, bots and workflow tools can each solve part of the job.

HIGH — **BUYER POWER**

Switching is easy before durable context/workflows form; buyers compare bundled and low-cost alternatives.

MED-HIGH — **SUPPLIER POWER**

Model, cloud, platform and connector providers affect cost, capability, policy and reliability.

MEDIUM — **NEW ENTRANTS**

LLM prototypes are easy; trustworthy multi-tenant operations, permissions, integrations and support are hard.

Net attractiveness: demand is real, but launch economics are only attractive when Aerealith owns a high-retention workflow, routes commodity tasks cheaply, and makes switching costly through customer-controlled context rather than lock-in.

- Strategic response: multi-model routing reduces supplier exposure; portable deployment reduces cloud dependence; open adapters reduce integration bottlenecks; workflows and context increase customer retention only after trust is established.

The launch should be managed as a trust-and-operations program

Likelihood and impact are pre-mitigation judgments for the recommended beachhead. Owners are functional accountabilities, not named individuals.

Risk	L	I	Leading indicator	Mitigation / owner
Permission or cross-workspace data leak	M	Critical	Unexpected source receipt; isolation test failure	Deny-by-default scopes, tenancy tests, red team / Security
Harmful or unauthorized action	M	Critical	Approval bypass, replay, duplicate execution	Risk classes, idempotency, approval, rollback / Product + Eng
Hallucinated or stale operational answer	H	High	Low citation coverage; stale source; corrections	Citations, freshness, abstention, eval set / AI Eng
COGS exceeds plan	H	High	AI COGS >20%; free MAU >$0.50	Routing, caps, BYOK, pricing / Finance + Product
Integration maintenance overload	H	High	Sync failures, support tickets, schema drift	Supported matrix, diagnostics, partner adapters / Platform
Incumbent bundles equivalent features	H	High	Win-loss cites existing suite	Own community workflow and control signature / Strategy
Weak activation or novelty churn	M-H	High	<30% activation; <40% week-4 workspace retention	Narrow template, guided proof, interview churn / Growth
Platform fee / policy change	M	Medium	Discord economics or rules change	Channel mix, compliant packaging, web team product / GTM
Privacy / regulatory claim failure	M	High	Unmapped data, deletion failures, unreviewed copy	Data map, counsel, claim review, evidence / Legal + Privacy
Parasocial / minor-safety incident	L-M	Critical	Dependency language, age-control bypass, complaints	AI disclosure, age-aware limits, escalation, NHP separation / Safety
Self-host support burden	M	High	Version fragmentation, insecure configs	Signed updates, supported matrix, telemetry opt-in, paid support / Platform

- Board-level triggers: any confirmed cross-tenant leak, unauthorized consequential action, or serious safety incident pauses expansion and initiates incident review before growth resumes.

North star: weekly active workspaces completing a trusted context-to-action loop

A loop counts only when the workspace retrieves permitted context with a citation and completes or explicitly declines a previewed action. Chat volume alone is not success.

WEEKLY TRUSTED LOOPS

Context retrieved + source inspected or accepted + action preview resolved + receipt recorded

Layer	Primary metric	6-month evidence target	12-month evidence target	Watch-out
Acquisition	Qualified workspace installs / signups	100 / month	500 / month	Unqualified bot traffic
Activation	2 sources + cited answer + resolved action in 7 days	>30%	>45%	Chat-only activation
Engagement	Weekly active activated workspaces	>45%	>55%	Notification-driven vanity use
Retention	Week-4 workspace retention	>40%	>50%	One-admin dependency
Trust	Cited-answer coverage; unauthorized actions	>90%; 0	>95%; 0	Low-risk actions masking failures
Economics	AI COGS / revenue; gross margin	<25%; >65%	<20%; >75%	Annual plans hiding negative cohorts
Expansion	Personal-to-workspace; community-to-team	Baseline	10%; 5%	Discount-driven upgrades
Reliability	Connector success; action success; MTTR	>98%; >97%; <4h	>99%; >99%; <1h	Retries hiding root cause

- Targets are management hypotheses. Reset them after beta cohorts establish baseline behavior; keep the metric definitions stable so improvements are real.

Test the riskiest assumptions before scaling build or acquisition

Each experiment has a disconfirming threshold. A failed test changes segment, workflow, package, or architecture rather than being explained away.

Hypothesis	Fastest credible test	Pass signal	Disconfirm / response
Community operators feel an urgent cross-tool problem	15 interviews + workflow diary from 5 communities	10 name weekly pain; 5 share data/process	<5 urgent -> narrow to another technical workflow
Permission visibility increases trust	Clickable prototype A/B with source and memory receipts	70% prefer receipt version; can explain scope	<50% -> redesign; terminology may be too technical
One cited brief creates repeat value	Concierge weekly brief across Discord/GitHub/docs	4/5 pilots request week 4; saves 1+ hour/week	<3 retain -> change job, cadence, or sources
Approved action is more valuable than answer-only	Prototype preview + manual execution behind scenes	60% of activated pilots approve 2+ actions/week	<30% -> action is premature or wrong target
Community buyers accept $49/month	Van Westendorp + paid pilot / refundable deposit	5 paid or deposit commitments from 10 qualified	<3 -> change unit, proof, scope, or price
Team expansion exists	Invite/admin workflow in retained communities	20% invite 3+ members; 2 request team controls	No expansion -> separate team acquisition motion
Margins can exceed 75%	Instrument every beta request and sync	Projected steady cohort COGS <25% revenue	>35% -> route/cap/reprice before GA

Research discipline: recruit from the intended beachhead, record current behavior and cost before showing the solution, separate enthusiasm from payment, and preserve failed findings. The goal is not to validate Aerealith; it is to find the version buyers repeatedly choose.

The next 90 days: decide the wedge, prove the loop, resolve the commercial contract

This sequence is designed to reduce strategic uncertainty before engineering scope expands.

WEEKS 1-2

- Appoint one beachhead owner
- Reconcile pricing and hosting documents
- Define permission/source schema
- Draft research screener and interview 5 buyers
- Instrument cost model by request type

Output
One strategy

WEEKS 3-4

- Complete 15-20 interviews
- Prototype receipts, memory controls and action preview
- Select one flagship Discord + GitHub workflow
- Recruit 5 design partners
- Define safety and incident stop criteria

Output
One prototype

WEEKS 5-8

- Run concierge / assisted pilots
- Measure baseline time and repeated questions
- Ship read-only context loop
- Simulate or manually fulfill approved action
- Publish internal weekly evidence review

Output
One retained workflow

WEEKS 9-12

- Charge or collect deposit from qualified pilots
- Finalize Free / Personal / Power / Community packaging
- Decide managed-cloud vs local-bridge launch promise
- Create trust packet and security roadmap
- Make go / pivot / stop gate decision

Output
One evidence-based decision

- Do not begin broad paid acquisition, a connector marketplace, or enterprise self-host commitments inside this window. Those choices become easier after the beachhead loop and trust contract are proven.

Continue, pivot, or stop based on observed behavior

Pre-commitment protects Aerealith from roadmap momentum. The thresholds below are strategic guardrails, not promises to investors or customers.

Gate	Continue	Pivot / Pause
DAY 30 **PROBLEM**	15+ qualified interviews; 10 report weekly pain; 5 commit to a pilot.	Pivot the segment or workflow if urgency is weak.
DAY 90 **PROOF**	4/5 pilots complete activation; 3 retain to week 4; 3 pay or deposit; no critical permission defect.	Rework onboarding/value loop; stop growth if trust fails.
MONTH 6 **PRODUCT**	30%+ activation; 40%+ week-4 workspace retention; repeat actions; projected gross margin >65%.	Reprice, cap, or narrow before public scale.
MONTH 12 **MOTION**	50%+ retained cohorts; positive contribution; repeatable acquisition; organic workspace expansion.	Hold team/enterprise expansion until PLG motion works.

- Hard stop criteria override growth metrics: confirmed cross-tenant data exposure, unauthorized consequential action, inability to honor deletion/retention commitments, or serious unmanaged user-safety harm.

The strategic answer: narrow the promise, deepen the control plane, expand by adjacency

Aerealith has a credible opportunity if it becomes the operating layer technical communities trust - and resists the temptation to launch the entire vision at once.

01 **Own a wedge**
Permission-aware AI operations for technical communities and developer-led organizations.

02 **Sell a loop**
Cited context -> visible memory -> preview -> approval -> action receipt.

03 **Simplify pricing**
Free, Personal, Power, Community, Team, and a quoted hybrid path.

04 **Instrument trust**
Treat permissions, provenance, correction, deletion and recovery as UI.

05 **Sequence integrations**
Discord, GitHub, docs/calendar, then read-only infrastructure; depth before breadth.

06 **Use model choice**
Make models replaceable and route by task; own the operating context above them.

07 **Protect margin**
Meter expensive usage openly; price platform fees and support into the channel.

08 **Prove retention**
Do not scale acquisition until the narrow workflow survives week 4.

09 **Separate research**
Keep NHP exploration governed and distinct from near-term commercial claims.

10 **Advance by gates**
Let evidence decide when Aerealith earns team, self-host, marketplace and enterprise scope.

- Recommended immediate decision: approve the technical-community beachhead and 90-day validation program; defer final public pricing and broad-market positioning until paid pilot evidence is available.

Model assumptions, scoring rules, and reproducibility notes

This page is the audit trail for the report's analytical judgments. Replace any assumption with better customer evidence; the formulas remain transparent.

MARKET SIZE

- TAM = 10.0M paid individuals x $120 + 250K communities x $600 + 500K teams x 10 seats x $240 + 10K hybrid enterprises x $50K = $3.05B ARR.
- SAM = 1.50M individuals x $120 + 50K communities x $600 + 75K teams x $2,400 + 2K enterprises x $35K = $460M ARR.
- SOM scenarios use explicit account counts and the same or lower annual revenue units. They are strategic scenarios, not probability-weighted forecasts.

COMPETITIVE SCORING

- 0 indicates no relevant documented capability found; 5 indicates category-leading public evidence. Scores are relative and directional.
- Vendor materials establish pricing and claimed/documented capability, not real-world product quality, reliability, adoption, or customer satisfaction.
- Aerealith is scored as target state from internal plans and is labeled accordingly. No planned feature is represented as currently shipped.

PRICING + ECONOMICS

- Public prices are snapshots as of 13 August 2026. Taxes, regional pricing, promotions, minimum seats, annual commitments, credits, and overages can change effective price.
- No currency conversion is used in competitor tables; native public prices are retained where relevant. Discord transaction processing is illustrated at 3%, explicitly as a planning assumption.
- Recommended Aerealith prices, plan-level COGS, gross-margin targets, churn, CAC, payback, five-year P&L; and funding envelope are management hypotheses to test.

EVIDENCE LIMITS

- Adoption sources use different samples, geographies and definitions; they indicate direction and maturity gaps rather than one harmonized time series.
- GitHub, Discord and small-business populations are demand and reach proxies. They do not independently validate addressable accounts or willingness to pay.
- Private-company revenue and retention are generally unavailable. No estimate in this report should be presented as audited financial or investment advice.

- Reproducibility: external facts map to S01-S58 and S62-S68. Internal assumptions map to S59-S61. Full URLs and evidence notes follow.

Source register 1 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S01	PUBLISHED / PRIMARY	Stanford HAI - 2025 AI Index: Economy Enterprise AI adoption in 2024; survey-based.	hai.stanford.edu/ai-index/2025-ai-index-report/economy
S02	PUBLISHED / PRIMARY	Stanford HAI - 2026 AI Index: Economy 2025 adoption and investment indicators.	hai.stanford.edu/ai-index/2026-ai-index-report/economy
S03	PUBLISHED / PRIMARY	McKinsey - The state of AI: Global survey 2025 AI adoption, scaling and agent experimentation.	www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai
S04	PUBLISHED / PRIMARY	Microsoft WorkLab - AI at Work Is Here. Now Comes the Hard Part Knowledge-worker use, BYOAI and leadership readiness.	www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part
S05	PUBLISHED / PRIMARY	Microsoft WorkLab - Agents, human agency, and the opportunity for every organization 2026 agentic-work framing; vendor-sponsored research.	www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization
S06	PUBLISHED / PRIMARY	U.S. Census Bureau - AI Use Among U.S. Businesses Business Trends and Outlook Survey estimates through 2026.	www.census.gov/library/stories/2026/05/ai-use-businesses.html
S07	PUBLISHED / PRIMARY	U.S. Census Bureau - Firm Adoption of Artificial Intelligence Working paper on firm- and employment-weighted adoption.	www.census.gov/library/working-papers/2026/adrm/CES-WP-26-25.html
S08	PUBLISHED / PRIMARY	OECD - AI use by individuals surges across the OECD Cross-country firm adoption trend.	www.oecd.org/en/about/news/announcements/2026/01/ai-use-by-individuals-surges-across-the-oecd-as-adoption-by-firms-continues-to-expand.html
S09	PUBLISHED / PRIMARY	OECD - Generative AI and the SME Workforce SME use patterns and limited integration into core work.	www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en/full-report/component-4.html

Source register 2 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S10	PUBLISHED / PRIMARY	Gartner - Worldwide AI Platforms and Models Forecast Category-spend context; not Aerealith TAM.	www.gartner.com/en/newsroom/press-releases/2026-07-20-gartner-forecasts-worldwide-ai-platforms-and-models-market-to-grow-63-percent-in-2026
S11	PUBLISHED / PRIMARY	Gartner - Worldwide AI Spending Forecast Broad AI-spend climate; dominated by infrastructure and vendor spend.	www.gartner.com/en/newsroom/press-releases/2026-05-19-gartner-forecasts-worldwide-ai-spending-to-grow-47-percent-in-2026
S12	PUBLISHED / PRIMARY	Gartner - Task-specific agents in enterprise applications Forecast for agent functionality in enterprise applications.	www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025
S13	PUBLISHED / PRIMARY	IDC - Global AI and Generative AI Spending Broader market-spend forecast; not targetable revenue.	www.idc.com/resource-center/blog/a-deep-dive-into-idcs-global-ai-and-generative-ai-spending/
S14	PUBLISHED / PRIMARY	McKinsey - State of AI trust in 2026 AI incidents and trust/readiness context.	www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/state-of-ai-trust-in-2026-shifting-to-the-agentic-era
S15	PUBLISHED / PRIMARY	Stack Overflow - 2025 Developer Survey: AI Developer trust and accuracy concerns.	survey.stackoverflow.co/2025/ai
S16	PUBLISHED / PRIMARY	GitHub - Octoverse 2025 Developer, repository and AI-project scale; vendor reporting.	github.blog/news-insights/octoverse/octoverse-a-new-developer-joins-github-every-second-as-ai-leads-typescript-to-1/
S17	PUBLISHED / PRIMARY	U.S. SBA Office of Advocacy - Frequently Asked Questions About Small Business 2026 U.S. small-business counts and employment.	advocacy.sba.gov/2026/02/03/frequently-asked-questions-about-small-business-2026/
S18	VENDOR SOURCE	Discord - About Discord Daily active user and audience profile figures; company reporting.	discord.com/company

Source register 3 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S19	VENDOR SOURCE	Discord - Discord appoints Humam Sakhnini as CEO Monthly active user and play-time figures; company reporting.	discord.com/press-releases/discord-appoints-new-ceo-humam-sakhnini
S20	VENDOR SOURCE	Discord - Discord Developer Platform App and automation platform capabilities.	discord.com/developers
S21	VENDOR SOURCE	Discord Developer Docs - Implementing App Subscriptions User and guild subscription mechanics.	docs.discord.com/developers/monetization/implementing-app-subscriptions
S22	VENDOR SOURCE	Discord Support - Premium Apps monetization support Platform fee structure and premium-app requirements.	support-dev.discord.com/hc/en-us/articles/23810643331735-Premium-Apps-Required-Support-for-Monetizing-Apps
S23	VENDOR SOURCE	Discord Support - Premium Apps payout Developer payout mechanics and fee tiers.	support-dev.discord.com/hc/en-us/articles/17299902720919-Premium-Apps-Payout
S24	VENDOR SOURCE	OpenAI - ChatGPT Business pricing Business workspace features and public pricing.	openai.com/business/pricing/
S25	VENDOR SOURCE	Anthropic - Claude Team plan Team plan public pricing and structure.	support.anthropic.com/en/articles/9266767-what-is-the-claude-team-plan
S26	VENDOR SOURCE	Anthropic - Claude Enterprise Enterprise controls and pricing structure.	www.anthropic.com/enterprise
S27	VENDOR SOURCE	Microsoft - Microsoft 365 Copilot pricing Public business pricing and bundle economics.	www.microsoft.com/en-us/microsoft-365-copilot/pricing

Source register 4 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S28	VENDOR SOURCE	Google - Google Workspace pricing Workspace bundles with Gemini features.	workspace.google.com/pricing
S29	VENDOR SOURCE	Notion - Notion pricing Business AI bundle and custom-agent credit model.	www.notion.com/pricing
S30	VENDOR SOURCE	Atlassian - Rovo for all paid cloud customers Rovo inclusion strategy and standalone plan.	www.atlassian.com/blog/announcements/team25-europe-rovo-everywhere
S31	VENDOR SOURCE	Atlassian - Rovo licensing Rovo Dev pricing and credit model.	www.atlassian.com/licensing/rovo
S32	VENDOR SOURCE	Atlassian - FY26 Q2 shareholder letter Rovo adoption and Teamwork Graph scale; company reporting.	www.atlassian.com/blog/announcements/shareholder-letter-q2fy26
S33	VENDOR SOURCE	Glean - Work AI for enterprise Permission-aware enterprise search and agent claims.	www.glean.com/
S34	VENDOR SOURCE	Glean Docs - Enterprise Flex pricing Per-user plus pooled-credit pricing model.	docs.glean.com/glean-enterprise-flex-pricing
S35	VENDOR SOURCE	Guru - Guru pricing Governed knowledge, integrations and usage-based structure.	www.getguru.com/pricing
S36	VENDOR SOURCE	Zapier - Zapier pricing Task-based plans and app-ecosystem scale.	zapier.com/pricing

Source register 5 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S37	VENDOR SOURCE	n8n - n8n pricing Execution-based cloud and self-hosted plans.	n8n.io/pricing/
S38	VENDOR SOURCE	Make - Make pricing Credit-based automation plans.	www.make.com/en/pricing
S39	VENDOR SOURCE	Dify - Dify pricing Hosted workspace plan and limits.	dify.ai/pricing
S40	VENDOR SOURCE	Dify - Dify Enterprise pricing Community edition, enterprise controls and private deployment.	dify.ai/pricing/dify-enterprise
S41	VENDOR SOURCE	Mem - Mem Personal AI memory positioning.	get.mem.ai/
S42	VENDOR SOURCE	Mem Docs - Mem MCP overview Interoperability with external AI clients.	docs.mem.ai/mcp/overview
S43	VENDOR SOURCE	Personal AI - Personal AI thesis Long-term memory and structured identity positioning.	www.personal.ai/thesis
S44	VENDOR SOURCE	Personal AI - Personal AI platform Programmable private AI and compliance claims.	www.personal.ai/platform
S45	VENDOR SOURCE	Limitless - Limitless Acquisition/wind-down precedent for personal memory hardware.	www.limitless.ai/

Source register 6 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S46	VENDOR SOURCE	MEE6 - MEE6 Premium Community-bot pricing and promotions.	mee6.xyz/en/premium
S47	VENDOR SOURCE	MEE6 Help - MEE6 services overview Separate Premium, AI and character subscriptions.	help.mee6.xyz/en/articles/611792-understanding-mee6-services-from-premium-to-ai-to-pro
S48	VENDOR SOURCE	Dyno - Dyno Modular moderation and server-management capabilities.	dyno.gg/
S49	VENDOR SOURCE	Carl-bot - Carl-bot Reaction roles, automod, logging and command features.	carl.gg/
S50	VENDOR SOURCE	Ticket Tool - Ticket Tool FAQs Community support-bot pricing.	docs.tickettool.xyz/debugging/faqs
S51	PUBLISHED / PRIMARY	European Commission - EU AI Act Application timeline and risk-based obligations.	digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai
S52	PUBLISHED / PRIMARY	NIST - AI Risk Management Framework and GenAI Profile Govern, Map, Measure and Manage risk framework.	www.nist.gov/itl/ai-risk-management-framework
S53	PUBLISHED / PRIMARY	U.S. FTC - Crackdown on deceptive AI claims Enforcement against unsupported AI and earnings claims.	www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes
S54	PUBLISHED / PRIMARY	U.S. FTC - Uphold privacy and confidentiality commitments Privacy promises and model-training representations.	www.ftc.gov/policy/advocacy-research/tech-at-ftc/2024/01/ai-companies-uphold-your-privacy-confidentiality-commitments

Source register 7 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S55	PUBLISHED / PRIMARY	UK ICO - Guidance on AI and data protection Data protection, fairness and lifecycle guidance.	ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/guidance-on-ai-and-data-protection/
S56	PUBLISHED / PRIMARY	UK ICO - ICO tech futures: Agentic AI Responsibility and novel data-protection risks in agentic systems.	ico.org.uk/about-the-ico/research-reports-impact-and-evaluation/research-and-reports/technology-and-innovation/tech-horizons-and-ico-tech-futures/ico-tech-futures-agentic-ai/
S57	VENDOR SOURCE	Discord - Getting global age assurance right Age-assurance timing, transparency and privacy context.	discord.com/blog/getting-global-age-assurance-right-what-we-got-wrong-and-whats-changing
S58	PUBLISHED / PRIMARY	Simple Icons - Simple Icons CC0 icon source used for comparative identification.	simpleicons.org/
S59	INTERNAL SOURCE	Aerealith AI - Aerealith AI Brand Guide v8 Internal brand and positioning source, supplied by the project owner.	Internal file supplied by Aerealith AI; no public URL.
S60	INTERNAL SOURCE	Aerealith AI - Aerealith AI Master Plan v0.3 Internal product, launch and integration plan, supplied by the project owner.	Internal file supplied by Aerealith AI; no public URL.
S61	INTERNAL SOURCE	Aerealith AI - Aerealith AI Pricing Model v3 Internal eight-tier pricing model, supplied by the project owner.	Internal file supplied by Aerealith AI; no public URL.
S62	VENDOR SOURCE	Anthropic - Claude API pricing Current model, caching, batch and tool-use pricing; cost benchmark, not a committed vendor choice.	platform.claude.com/docs/en/about-claude/pricing
S63	VENDOR SOURCE	Google AI for Developers - Gemini Developer API pricing Current token, caching, batch and grounding prices; several rates change in 2027.	ai.google.dev/gemini-api/docs/pricing

Source register 8 of 8

Direct links are provided for external evidence. Vendor pages are identified so feature and pricing statements are not mistaken for independent validation.

S64	VENDOR SOURCE	Stripe - Pricing and fees Public payment-processing price reference; geography, payment method and negotiated terms vary.	stripe.com/pricing
S65	VENDOR SOURCE	Cloudflare Developers - Workers pricing Published request, CPU and base-plan economics for an edge-compute reference architecture.	developers.cloudflare.com/workers/platform/pricing/
S66	VENDOR SOURCE	Supabase - Pricing and fees Published database, MAU, storage, egress and compute allowances used as infrastructure reference points.	supabase.com/pricing
S67	PUBLISHED / PRIMARY	Benchmarkit - 2025 SaaS Performance Metrics Survey benchmark: 77% median total gross margin, 81% subscription gross margin and 30% professional-services margin.	www.benchmarkit.ai/2025benchmarks
S68	PUBLISHED / PRIMARY	SaaS Capital - 2026 bootstrapped SaaS benchmarks Survey of 1,000+ private B2B SaaS companies; 103% median NRR and 91% median GRR in the stated cohort.	www.saas-capital.com/blog-posts/benchmarking-metrics-for-bootstrapped-saas-companies/

Brand, trademark, and use notices

The report uses Aerealith's supplied identity and third-party marks only to make competitive comparisons legible.

AEREALITH AI

A companion for a more connected world.

DOCUMENT CONTROL

Document	Aerealith AI Market Intelligence & Launch Strategy
Version	1.1 - expanded financial model
Evidence date	13 August 2026
Market model	Bottom-up; modeled estimates explicitly labeled
Confidentiality	Confidential strategic planning document

Aerealith materials

Brand Guide v8, Master Plan v0.3, and Pricing Model v3 were supplied by the project owner and treated as internal planning inputs.

Third-party marks

Logos and product names are used for comparative identification. All trademarks and brand rights belong to their respective owners. Inclusion does not imply endorsement, partnership, or sponsorship.

Icon source

Several monochrome product marks are rendered from the Simple Icons project. Simple Icons assets are CC0; its disclaimer notes that trademark rights remain with owners.

Advisory boundary

This report is strategic market research, not legal, tax, accounting, investment, security-certification, or regulatory advice. Verify current pricing, policies, and laws before making commitments.